

PACT-2014 The 23<sup>rd</sup> International Conference on  
Parallel Architectures and Compilation Techniques

# CAWS: Criticality-Aware Warp Scheduler for GPGPU Workloads

Shin-Ying Lee and Carole-Jean Wu

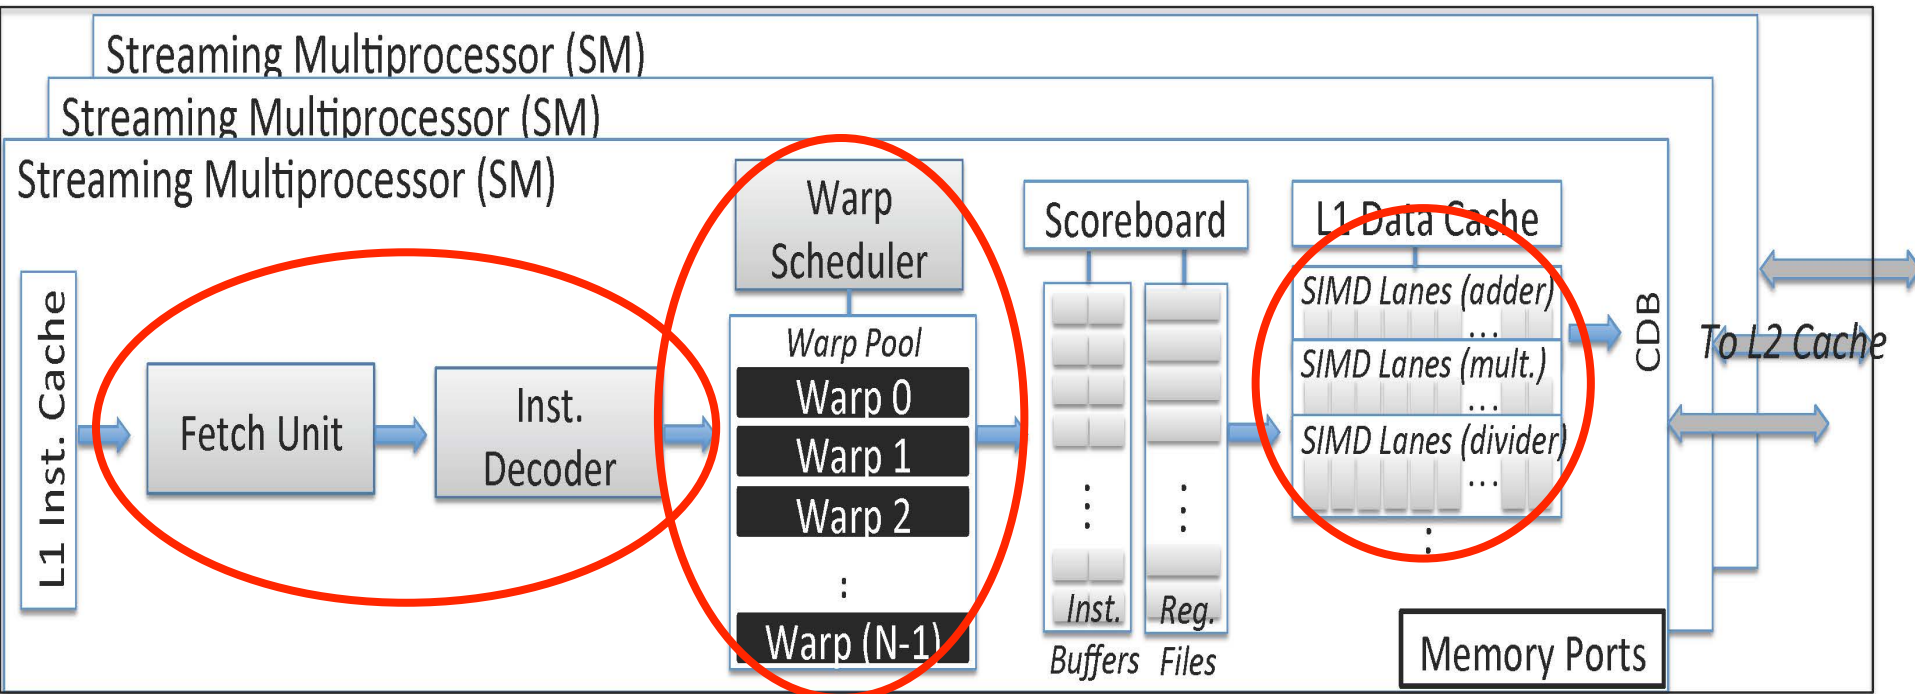
School of Computing, Informatics, and Decision Systems Engineering  
Arizona State University

# Graphics Processing Units (GPUs)

- In addition to render video frames, modern GPUs are also deployed to process parallel workloads.
- The benefit of using GPUs to perform general computation is to hide operation latency
  - *Massive multi-threading*
  - *Fast context-switching*

# Unified GPU Architecture (Computation Path)

- The computation and graphics unified GPU architecture to support general computation



# Research Questions

- Is the current GPU design good enough to hide execution latency?
- What kind of execution latency is hidden under the GPU scheduler?
- How do we propose a better scheduling policy to improve GPU's latency hiding ability?

# Outline

- GPU Latency Characterization
- Warp/Wavefront Criticality
- Criticality-Aware Warp Scheduling (CAWS)
- Discussion and Conclusion

# Root-cause of Warp Stall

- Pipeline hazards
  - Data hazard
  - Structural hazard
  - Control hazard
- Memory access latency
- Instruction fetch policy
- Synchronization barrier
- Warp scheduling policy

# Latency Characterization Algorithm

*probe(w)*

*w.Scheduling* += *CurTime* - *w.PrevTime*

*w.PrevTime* = *CurTime*

*if*(*Sync*)

*w.Sync* ++

*else if*(*EmptyInstBuffer*)

*w.Fetch* ++

*else if*(*BranchTaken*)

*w.CtrlHazard* ++

*else if*(*DataDependency*)

*if*(*OldDataCacheMiss*)

*w.DataCacheMiss* ++

*else*

*w.DataHazard* ++

*else if*(*FU\_unavailable*)

*w.StructuralHazard*

*else*

*w.Exec* ++


# Warp/Wavefront Scheduler

1. Select a warp according to the scheduling policy at every cycle
2. Probe and record the selected warp's current status
3. Iteratively till find a ready warp

```
while(Not VisitedAllWarps)  
    w = FindNextWarp()  
    probe(w)  
    if(w is a ready warp)  
        issue(w)  
        break
```



# Latency Characterization Results

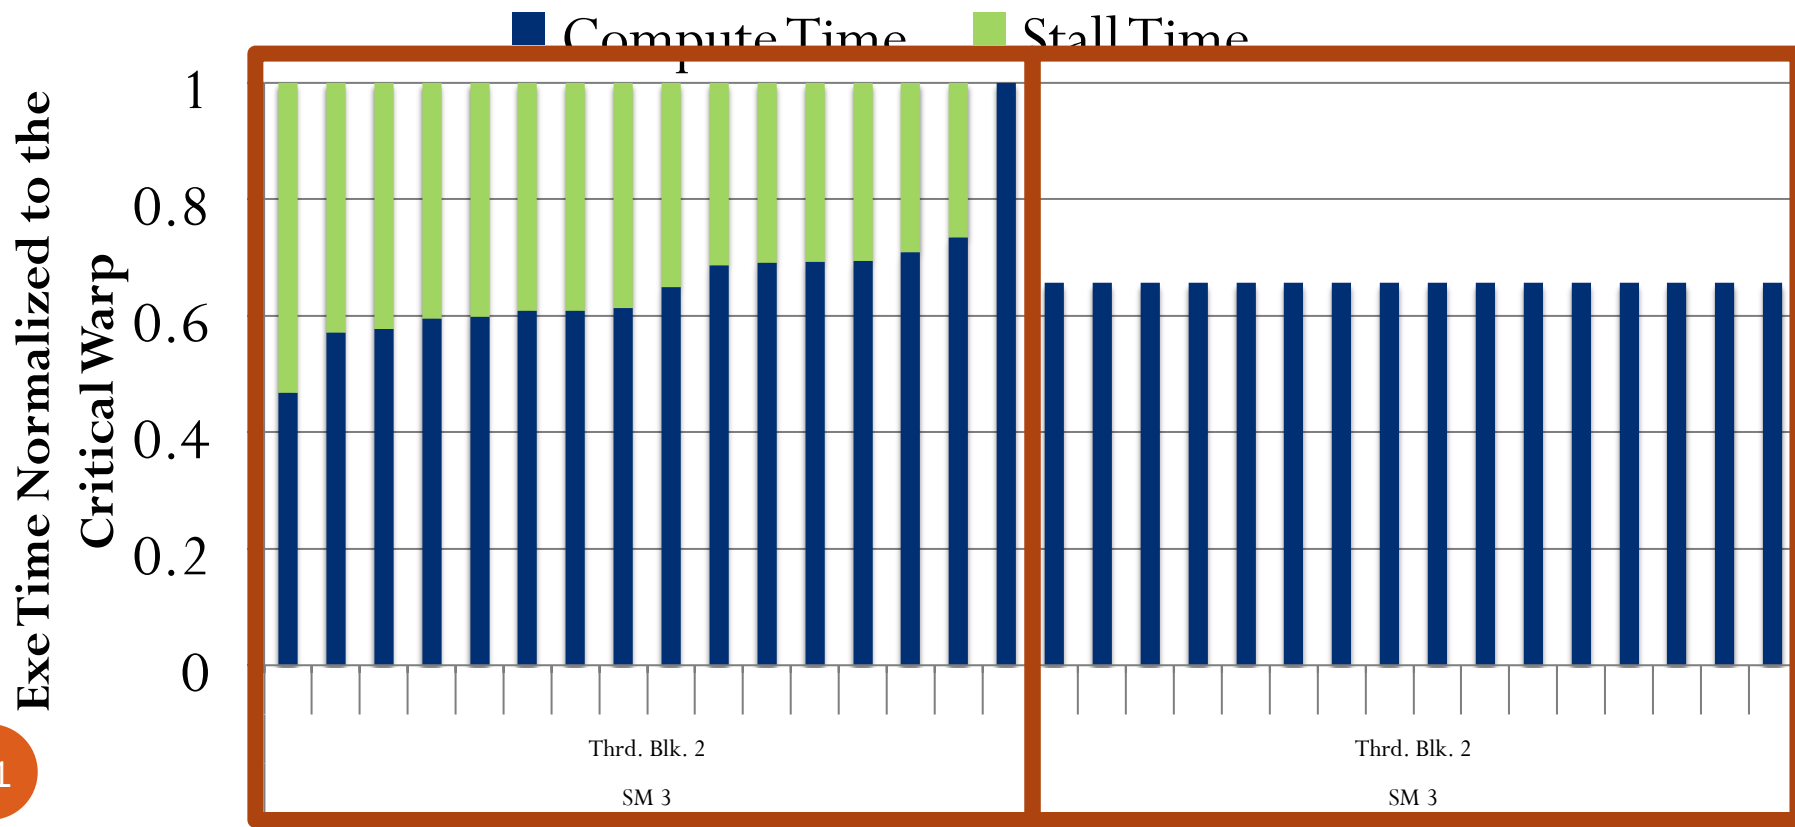


# Outline

- GPU Latency Characterization
- Warp/Wavefront Criticality
- Criticality-Aware Warp Scheduling (CAWS)
- Discussion and Conclusion


# Warp/Wavefront Criticality

- Slow warps/wavefronts take much longer time to finish their designated jobs.
- Fast warps are blocked at an implicit/explicit synchronization barrier to wait for the slow warps/wavefronts



# Factor Causing Warp Criticality for *bfs*

**workload imbalance caused by branches**



# Outline

- GPU Latency Characterization
- Warp/Wavefront Criticality
- Criticality-Aware Warp Scheduling (CAWS)
- Discussion and Conclusion

# Criticality-Aware Warp Scheduling

- To equalize the execution time disparity, the Criticality-Aware Warp Scheduling (CAWS) algorithm prioritizes and schedule warps by warps' weight.
- Assign warps different weight based on their criticality.
- Slower warps receive more time slots to run advance.

# Types of CAWS Policies

- CAWS-Blk
  - Prioritizing warps within a thread block
  - To equalize execution time disparity within a thread block
- CAWS-SM
  - Prioritizing warps within as SM
  - To equalize execution time disparity within an SM
- CAWS-Avg
  - Identifying the critical warps based on the average execution time across an SM

# Experimental Environment

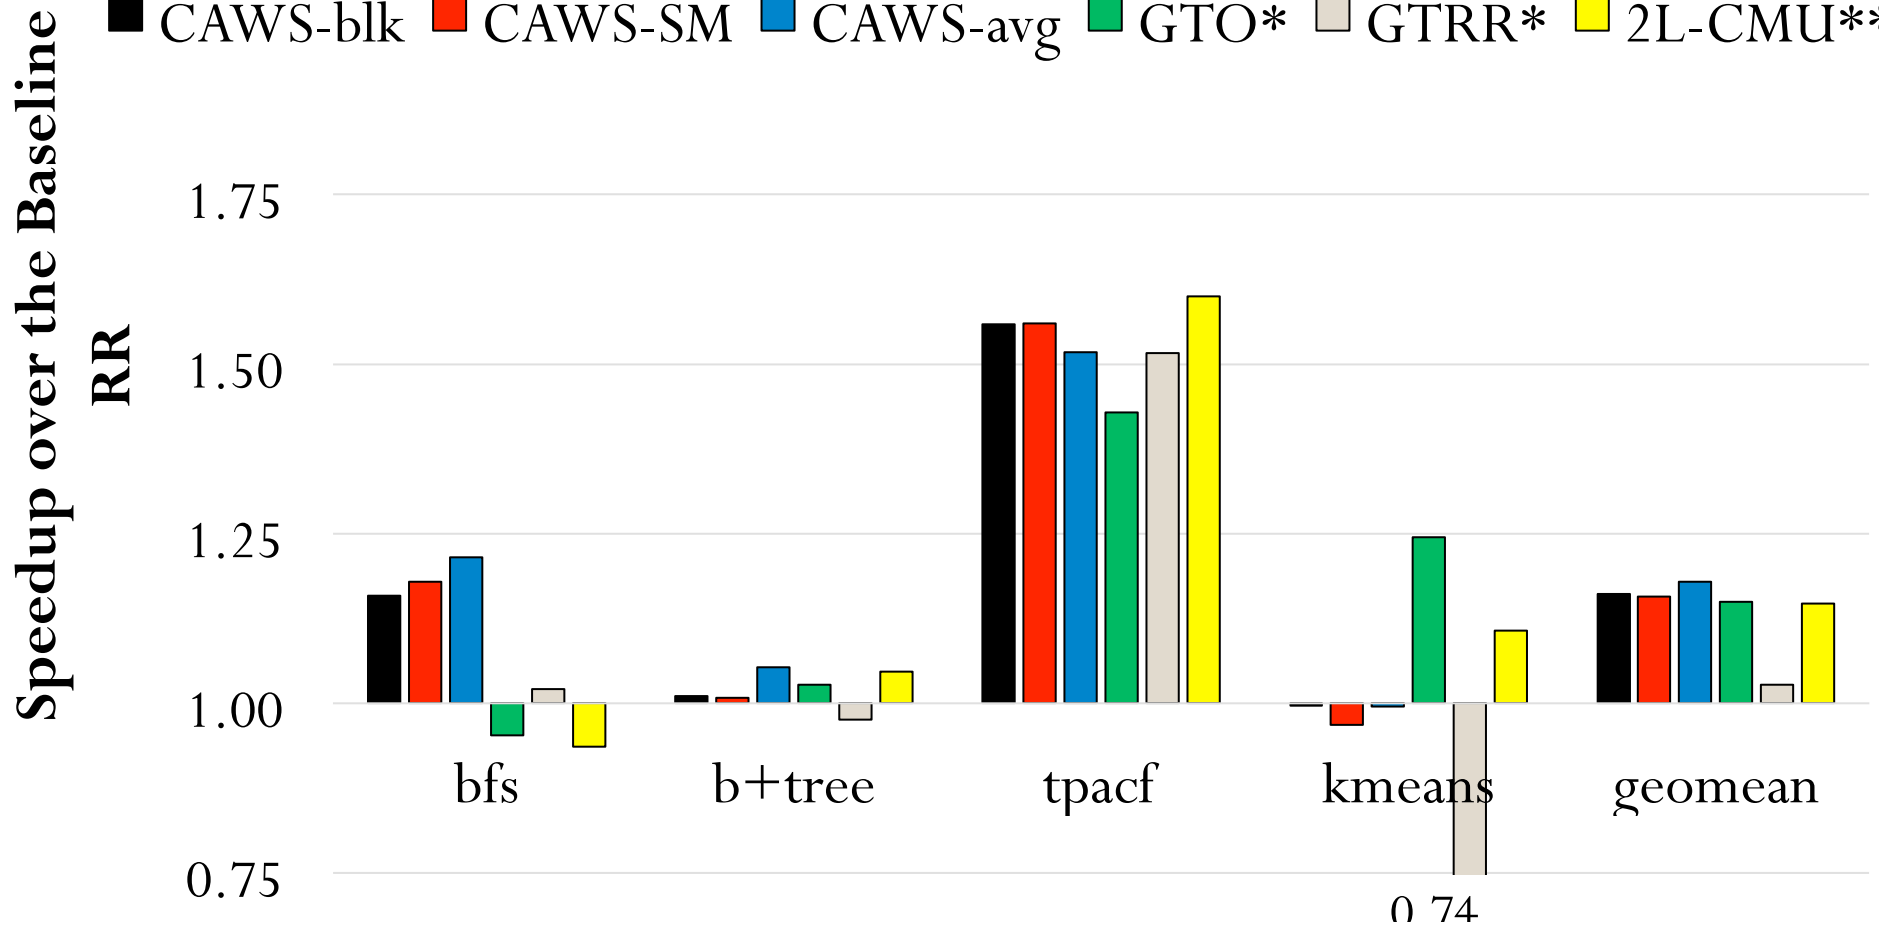
- GPU simulation infrastructure
  - GPGPU-sim v3.2.0\*
  - nVIDIA nvcc toolchain v3.2
  - nVIDIA GTX480 architecture
- Benchmarks
  - Imbalance workloads: *bfs*
  - Small parallel regions: *b+tree*
  - I-cache intensive: *tpacf*
  - D-cache intensive: *kmeans*

\*A. Bakhoda, et al, “*Analyzing CUDA Workloads Using a Detailed GPU Simulator*,” ISPASS-2009, 2009



# CAWS Speedup with Oracle Knowledge


■ CAWS-blk ■ CAWS-SM ■ CAWS-avg ■ GTO\* ■ GTRR\* ■ 2L-CMU\*\*



\*T. G. Rogers, "Cache-conscious Wavefront Scheduling," MICRO-45, 2012

\*\*V. Narasiman, et al. "Improving GPU Performance via Large Warps and Two-level Scheduler," MICRO-44, 2011

# Latency Characterization of *bfs*



# CAWS with Criticality Prediction on *bfs*

Speedup Normalized to Baseline

RR Scheduler

1.25  
1.20  
1.15  
1.10  
1.05  
1.00  
0.95  
0.90

Round-Robin (RR)

CAWS-blk

CAWS-SM

CAWS-avg

CAWS-blk

CAWS-SM

CAWS-avg

Oracle

Criticality-Guided Predictor

# Outline

- GPU Latency Characterization
- Warp/Wavefront Criticality
- Criticality-Aware Warp Scheduling (CAWS)
- Discussion and Conclusion

# Research Questions

- Is the current GPU design good enough to hide execution latency?
- What kind of execution latency is hidden under the GPU scheduler?
- How do we propose a better scheduling policy to improve GPU's latency hiding ability?

# Conclusion

- This is the first latency characterization algorithm for GPUs that enables a deep understanding of how well the latency hiding ability is for modern GPU schedulers
- We present the GPGPU workloads' results to indicate places for performance improvement.
- We design a family of CAWS policies to improve GPGPU workload performance by equalizing warp execution time.
- The CAWS policies can potentially achieve 17% of performance improvement on average.

# Thank You!

---

*CAWS: Criticality-Aware Warp Scheduler for GPGPU Workloads*

*Shin-Ying Lee and Carole-Jean Wu*

*Arizona State University*

**BACKUP**



# bfs Algorithm

while(not visited all node)

    kernel\_1

        travers all nodes' children in current depth

    /\* an implicit barrier here \*/

    kernel\_2

        go to next depth

    /\* an implicit barrier here \*/

# CAWS Implementation

- Implemented by count-down counters
  - Every warp has its own priority counter.
  - A priority counter's initial value is corresponding the warp's priority/criticality
  - The counter value decrements cycle by cycle.
  - Scheduler picks up the warp having the lowest counter value to be issued.
  - Once a warp is issued, its counter value is reset to the initial value.

# Criticality Inversion

- A fast warp becomes a new critical warp due to starving.
- It may limit the overall speedup or make performance even worse.

| Policy   | Speedup | Criticality Inversion |
|----------|---------|-----------------------|
| CAWS-blk | 1.16    | 8.89%                 |
| CAWS-SM  | 1.18    | 2.22%                 |
| CAWS-Avg | 1.21    | 0%                    |

# Comparison of CAWS Policies

|                 | Pros   | Cons  |
|-----------------|--|---|
| <b>CAWS-Blk</b> | To quickly finish a thread block   | Useless for thread blocks having no critical warp |
| <b>CAWS-SM</b>  | Useful for cases such that critical warps mapped to thee same thread block | Some thread blocks might get starving             |
| <b>CAWS-Avg</b> | Criticality inversion avoidance  | More complicated to implement                     |

# Latency Characterization Comparison

